

Computer Vision

Edge Detection

Augusto Silva

Department of Electronics, Telecommunications and Informatics
University of Aveiro

`augusto.silva@ua.pt`

Outline


- 1 Introductory Concepts
- 2 Mathematical Tools
- 3 Edge Detectors
- 4 Lines and Corners

Edge detectors

- Very important local image pre-processing methods used to locate changes in the intensity function;
- Edges play a key role in human and machine perception
- Sometimes only edge elements with strong magnitude (edgels) suffice for image understanding.
- Significant data reduction for image representation
- A change of the image function can be described by a gradient that points in the direction of the largest growth of the image function.
- An edge is a property attached to an individual pixel and is calculated from the image function behavior in a neighborhood of that pixel.

Edge Sources


- Problem: How to retrieve the relevant edges?



Adapted from Image Processing , Analysis, and Machine Vision 3Ed, Sonka et al

Edge Concepts


- An edge is a vector variable with two components, magnitude and direction
 - The edge magnitude is the magnitude of the gradient
 - The edge direction f is rotated with respect to the gradient direction by 90° .
 - The gradient direction points to the maximum growth of the function.
 - Ideally boundaries are made of pixels with edge magnitudes



Adapted from Image Processing , Analysis, and Machine Vision 3Ed, Sonka et al

Edge Profiles

- The edge profile is a plot of image values along the gradient direction (perpendicular to the edge direction).
- Typical profiles:



Adapted from Image Processing , Analysis, and Machine Vision 3Ed, Sonka et al

- Roof edges are typical for objects corresponding to thin lines in the image.
- Edge detectors are usually tuned for some type of edge profile.

Mathematical Tools

- The Gradient (continuous version) is a vector field requiring magnitude and angular descriptors

$$\nabla f(x, y) = \left(\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right) = (g_x, g_y)$$

$$|\nabla f(x, y)| = \sqrt{g_x^2 + g_y^2}, \quad \phi = \arg(g_x, g_y)$$

- Symmetrical discrete approximation:
 $\Delta_i g(i, j) = g(i, j) - g(i - n, j), \quad \Delta_j g(i, j) = g(i, j) - g(i, j - n)$
- Sometimes only edge magnitudes without regard to their orientation is important. The Laplacian is the linear differential operator of choice since it is rotationally invariant.

$$\nabla^2 f(x, y) = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

Discrete approximations

- Assymmetrical 1st derivative discrete approximation (including pixel (i, j)):

$$\begin{cases} \Delta_i g(i, j) &= g(i, j) - g(i - n, j) \\ \Delta_j g(i, j) &= g(i, j) - g(i, j - n) \end{cases}$$

- Symmetrical 1st derivative discrete approximation (excluding pixel (i, j)):

$$\begin{cases} \Delta_i g(i, j) &= g(i + n, j) - g(i - n, j) \\ \Delta_j g(i, j) &= g(i, j + n) - g(i, j - n) \end{cases}$$

- Possible Laplacian discretization

$$\nabla^2 f_{i,j} = (f_{i+1,j} + f_{i-1,j} + f_{i,j+1} + f_{i,j-1}) - f_{i,j}$$

Example


Exercise

- Read the image *stairs.tiff*
- Use Matlab functions `gradient` and `del2` to compute the gradient and laplacian of image
- Display the resulting images
- Superimpose a *quiver* plot to see the gradient directions


Original



Gradient



Laplacian



Edge sharpening

Exercise

- Compute the image Laplacian (use the *stairs.tiff* image)
- Observe the edge sharpening effects for various values of $C \geq 0$ when computing

$$f_{i,j} = f_{i,j} - C \times L_{i,j}$$

- Use a gaussian filter to compute a blurred version B of the image
- Observe the edge sharpening effects for various values of $C \geq 0$ when computing

$$f_{i,j} = f_{i,j} + C \times (f_{i,j} - B_{i,j})$$

- This operation is a digital version of the so called unsharp mask operation

Differential operators for edge detection

- Operators approximating derivatives of the image function using differences:
 - Rotationally invariant
 - Laplacian
 - First derivative approximation
 - Multi-mask
 - Orientation determined upon a best match criteria
- Operators based on the zero-crossings of the image function second derivative
 - Marr-Hildreth
 - Canny
- Operators which attempt to match an image function to a parametric model of edges.

Edges with first derivatives

- Individual gradient operators that examine small local neighborhoods are convolutions
- Operators which are able to detect edge direction are represented by a collection of masks, each corresponding to a certain direction.
- Roberts operator

$$h_1 = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix} \quad h_2 = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$$

$$|f(i, j)| = |g(i, j) - g(i + 1, j + 1)| + |g(i, j + 1) - g(i + 1, j)|$$

- Simple, fast but very noise sensitive

Compass operators

- The gradient is estimated in eight (for a 3×3 convolution mask) possible directions
- The convolution result of greatest magnitude indicates the gradient direction.
- Prewitt

$$h_1 = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix} \quad h_2 = \begin{bmatrix} 0 & 1 & 1 \\ -1 & 0 & 1 \\ -1 & -1 & 0 \end{bmatrix} \quad h_3 = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$$

- Sobel

$$h_1 = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} \quad h_2 = \begin{bmatrix} 0 & 1 & 2 \\ -1 & 0 & 1 \\ -2 & -1 & 0 \end{bmatrix} \quad h_3 = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

Compass operators

- Robinson

$$h_1 = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} \quad h_2 = \begin{bmatrix} 0 & 1 & 21 \\ -1 & 0 & 1 \\ -2 & -1 & 0 \end{bmatrix} \quad h_3 = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

- Kirsch

$$h_1 = \begin{bmatrix} 3 & 3 & 3 \\ 3 & 0 & 3 \\ -5 & -5 & -5 \end{bmatrix} \quad h_2 = \begin{bmatrix} 3 & 3 & 3 \\ -5 & 0 & 3 \\ -5 & -5 & 3 \end{bmatrix} \quad h_3 = \begin{bmatrix} -5 & 3 & 3 \\ -5 & 0 & 3 \\ -5 & 3 & 3 \end{bmatrix}$$

Vertical and horizontal edges

Exercise

- Use the image *stairs.tiff*.
- Obtain a horizontal edge function E_h with the Sobel operator h_1
- Obtain a vertical edge function E_v with the Sobel operator h_2
- Obtain a global edge function making

$$E_g = |E_h| + |E_v|$$

- Threshold each filtered image with the IP toolbox functions `graythresh` and finally obtain the binary edge map with `im2bw`
- Repeat the exercise with various instances of zero mean gaussian noise superimposed on the image.

The Laplacian

- Second derivative approximation of ∇^2
- Convolution mask implementation

4-neighborhood

$$h = \begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

8-neighborhood

$$h = \begin{bmatrix} 1 & 1 & 1 \\ 1 & -8 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

- Non rotationally invariant alternatives

$$h = \begin{bmatrix} 2 & 1 & 2 \\ 1 & -4 & 1 \\ 2 & 1 & 2 \end{bmatrix}$$

$$h = \begin{bmatrix} -1 & 2 & -1 \\ 2 & -4 & 2 \\ -1 & 2 & -1 \end{bmatrix}$$

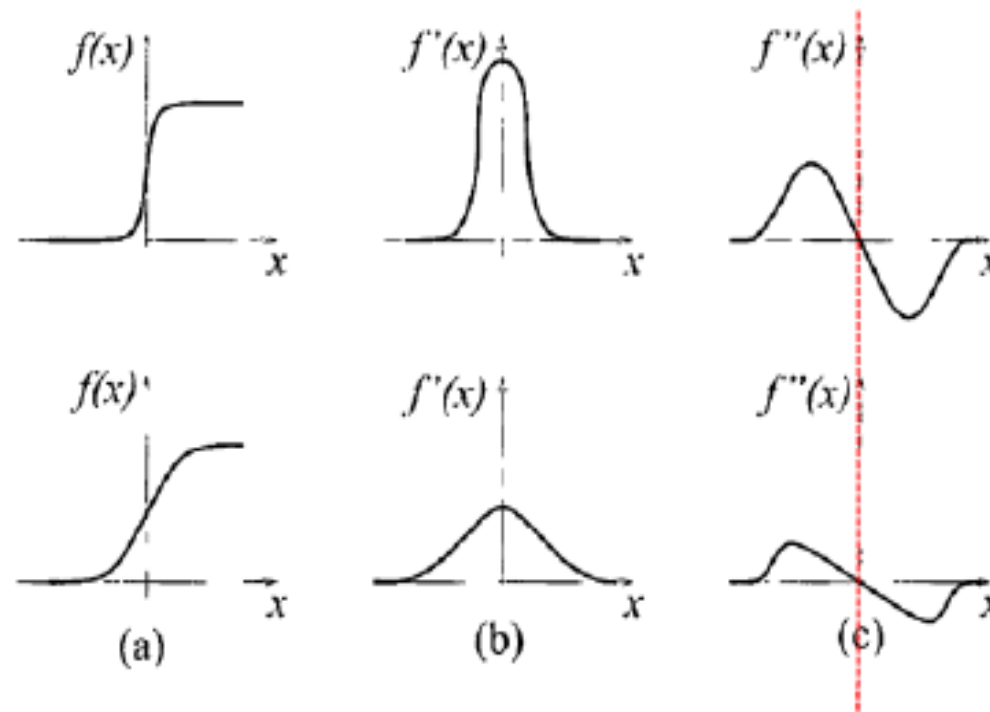
- Problem with Laplacians: double responses in some edge instances

Zero-crossings of the second derivative

- Marr's theory
- Neurophysiological experiments indicate that object boundaries are the most important cues that link an intensity image with its interpretation.
- Edge detection with first derivative approximations over small neighbourhoods are very noise sensitive and object dependent
- An edge detection technique based on the zero-crossings of the second derivative explores the fact that a step edge corresponds to an abrupt change in the image function.
 - The first derivative of the image function should have an extremum at the position corresponding to the edge
 - It is much easier and more precise to find a zero-crossing position than an extremum.

Edge profiles revisited

- Consistency for various edges profiles



Adapted from Image Processing , Analysis, and Machine Vision 3Ed, Sonka et al

- How to compute the second derivative robustly?

Laplacian of Gaussian - LoG

- The LoG operator integrates
 - A low-pass smoothing operator to minimize noise effects and control the influence of neighbour pixels
 - The Laplacian operator to compute the second derivative
- Continuous model:

$$\text{LoG} [f(x, y)] = \nabla^2 [G(x, y, \sigma) * * f(x, y)]$$

- The order of performing differentiation and convolution can be interchanged because of the linearity of the operators involved

$$\text{LoG} [f(x, y)] = [\nabla^2 G(x, y, \sigma)] * * f(x, y)$$

- Gaussian is rotationally invariant so it is useful to make

$$r^2 = x^2 + y^2, \quad G(x, y, \sigma) = G(r, \sigma) = e^{-r^2/\sigma^2}$$

The Mexican Hat function


Exercise

- Compute the second derivative $\nabla^2 G(r, \sigma)$
- Show that returning to the original co-ordinates x, y and introducing a normalizing multiplicative coefficient c , we get an inverted Mexican Hat function
- Make a 3D plot

$$h(x, y) = c \left(\frac{x^2 + y^2 - \sigma^2}{\sigma^4} \right) e^{-(x^2 + y^2)/2\sigma^2}$$

(Reversed) Kernel Approximation

$$\begin{bmatrix} 0 & 0 & -1 & 0 & 0 \\ 0 & -1 & -2 & -1 & 0 \\ -1 & -2 & 16 & -2 & -1 \\ 0 & -1 & -2 & -1 & 0 \\ 0 & 0 & -1 & 0 & 0 \end{bmatrix}$$



LoG issues

- Gaussian smoothing effectively suppresses the influence of the pixels that are more than a distance 3σ from the current pixel
 - A significant area surrounding the current pixel is taken into account; the influence of more distant points decreases according to the 3σ of the Gaussian.
- The Laplace operator is an efficient and stable measure of changes in the image.
- Convolution masks become large for larger σ for example, a $\sigma = 4$ needs a mask about 40 pixels wide.
- Separable decomposition of $\nabla^2 G$ exists to speed up the filtering process
- Be aware of sharp corner missing and fake close loops of edges

DoG: Difference of Gaussians

- The practical implication of Gaussian smoothing is that edges are found reliably.
 - Increase the value of σ if only globally significant edges are required
- The $\nabla^2 G$ operator can be very effectively approximated by a convolution with a mask that is the difference of two Gaussian averaging masks with substantially different σ
 - Difference of Gaussians (DoG) method
 - The ratio of standard deviations is an issue
- Coarse implementation $\nabla^2 G$
 - Filter the image twice with two different smoothing gaussian masks and obtain the difference

Looking for Zero Crossings

- Brute force zero search fails
- Naive search over interval produce piecewise disconnected edges
- Necessary to have a robust zero-crossing detector
 - Scan the LoG or DoG image with a 2x2 moving window
 - Assign an edge label to any one corner pixel, say the upper left, if LoG/DoG image values of both polarities occur in the 2x2 window.
 - No edge label would be given if values within the window are either all positive or all negative.
 - Use first derivative (gradient support)
 - Only those zero-crossings for which there is sufficient edge evidence from a first-derivative edge detector.

Canny Edge Detection

- Canny proposed an approach to edge detection that is optimal for step edges corrupted by white noise
- Criteria
 - The detection criterion expresses the fact that important edges should not be missed and that there should be no spurious responses.
 - The localization criterion says that the distance between the actual and located position of the edge should be minimal.
 - The one response criterion minimizes multiple responses to a single edge.

Canny Edge Detection: Fundamentals

- A step edge is given by its position, orientation, and possibly magnitude (strength).
- The driving ideas:
 - Convolve the image with a symmetric 2D Gaussian filter
 - Differentiate in the direction of the gradient (perpendicular to the edge direction) responses.
 - Note that LoG/DoG zero-crossing operator does not give information about edge direction, as it uses a Laplacian filter).
- A directional derivative of the Gaussian filter is given by

$$G_n = \frac{\partial G}{\partial \mathbf{n}} = \mathbf{n} \nabla G$$

- The direction \mathbf{n} should be oriented perpendicular to the edge

Canny Edge Detection: Edge estimation

- A good estimate of the perpendicular direction to the edge can be obtained by smoothed gradient direction

$$\mathbf{n} = \frac{\nabla (G * f)}{|(G * f)|}$$

- The edge location corresponds to the local maximum of the image f convolved with the operator $G_{\mathbf{n}}$ in the direction \mathbf{n}

$$G_n = \frac{\partial (G * f)}{\partial \mathbf{n}} = 0 \therefore \frac{\partial^2 (G * f)}{\partial \mathbf{n}^2} = 0$$

- Spurious responses to the single edge caused by noise usually create a 'streaking' problem
- Canny's procedure includes threshold with hysteresis

Canny Edge Detection: Multi-scale edge prediction

- Multi-scale prediction of edge strength
 - What's the correct scale of the operation?
 - Not known a priori
 - Canny proposed a feature synthesis approach.
 - Start by marking all significant edges from the operator with the smallest scale
 - Synthesize edges of a hypothetical operator with larger σ
 - Compare the edge response to the actual response for larger σ
 - Additional edges are marked only if they have a significantly stronger response than that predicted from synthetic output.
 - Repeat for various scales
 - Build a cumulative edge map

Canny Edge Detection: Algorithm

- ① Convolve an image f with a Gaussian of scale σ .
- ② Estimate local edge normal directions \mathbf{n} for each pixel in the image
- ③ Find the location of the edges using non-maximal suppression
- ④ Compute the magnitude of the edge
- ⑤ Threshold edges in the image with hysteresis to eliminate spurious responses.
- ⑥ Repeat steps 1) through 5) for ascending values of the standard deviation σ .
- ⑦ Aggregate the final information about edges at multiple scale using the 'feature synthesis' approach.

Experiments

Exercise

- Study the IP toolbox function `edge`
- Read the images *winchester.tiff* and *pentagon.tiff*
- Perform edge detection with
 - First derivative operators
 - Zero-crossing LoG operators
 - Canny operator with several values of σ .
 - Repeat for several levels of zero mean gaussian noise.

Line Detection


- Same rationale of detecting "roof" like profiles along "strategic" orientations: $0^\circ, 45^\circ, 90^\circ, 135^\circ \dots$
- Convolution Kernels

$$h_1 = \begin{bmatrix} -1 & -1 & -1 \\ 2 & 2 & 2 \\ -1 & -1 & -1 \end{bmatrix} \quad h_2 = \begin{bmatrix} 2 & -1 & -1 \\ -1 & 2 & -1 \\ -1 & -1 & 2 \end{bmatrix} \quad h_3 = \begin{bmatrix} -1 & 2 & -1 \\ -1 & 2 & -1 \\ -1 & 2 & -1 \end{bmatrix}$$

- Lines detected this way are collections of edges. Most of the time non single pixel wide edges.
- Necessary to introduce line thinning algorithms (ahead in this course)

Corner Detection

- Corners in images can be located using local detectors;
 - Input to the corner detector is the gray-level image
 - Output is the image in which values are proportional to the likelihood that the pixel is a corner. Interest points are obtained by thresholding the result of the corner detector.
- Edge detectors themselves are not stable at corners.
 - Gradient at the tip is ambiguous



Adapted from Image Processing , Analysis, and Machine Vision 3Ed, Sonka et al

- The corner in the image can also be defined as a pixel in its small neighborhood where there are two dominant and different edge directions. (Not very precise definition)

Corner Detectors

- Moravec Detector:
 - Auto-Correlation based

$$\text{MO}(i, j) = \frac{1}{8} \sum_{k=i-1}^{i+1} \sum_{l=j-1}^{j+1} |f(k, l) - f(i, j)|$$

- Compare the MO discriminant according to a predefined threshold
- Zuniga-Haralick:
 - Approximate the function $f(i, j)$ by a 2D cubic polynomial

$$f(i, j) = c_1 + c_2x + c_3y + c_4x^2 + c_5xy + c_6y^2 + c_7x^3 + c_8x^2y + c_9xy^2 + c_{10}y^3$$

- Consider the rate of change in gradient direction and build a discriminant $\text{ZH}(i, j)$

$$\text{ZH}(i, j) = \frac{-2(c_2^2c_6 - c_2c_3c_5 + c_3^2c_4)}{(c_2^2 + c_3^2)^{3/2}}$$

Harris corner detector

- Auto-correlation based
- Improvement upon Moravec's corner detector
- Use a sliding window W patch and estimate the sum of square differences of the discriminant function

$$S_w(\Delta x, \Delta y) = \sum_{x \in W} \sum_{y \in W} (f(x_i, y_j) - f(x_i - \Delta x, y_j - \Delta y))^2$$

- A corner point must have a high response of $S_w(\Delta x, \Delta y)$ for all $\Delta x, \Delta y$.
- With a first order Taylor approximation near (i, j)

$$f(x_i - \Delta x, y_j - \Delta y) \approx f(x_i, y_j) + \left[\frac{\partial f(x_i, y_j)}{\partial x}, \frac{\partial f(x_i, y_j)}{\partial y} \right] \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix}$$

Harris corner detector

- After some algebra the analytical maximum of $S_w(\Delta x, \Delta y)$ is given by

$$[\Delta x, \Delta y] A_W(x_i, y_j) \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix}$$


- Where $A_W(x_i, y_j)$ denotes the Harris matrix containing the second derivatives around the central point of S_w

$$A(x_i, y_j) = \begin{bmatrix} \sum_{x_i, y_j \in W} \frac{\partial^2 f(x_i, y_j)}{\partial x^2} & \sum_{x_i, y_j \in W} \frac{\partial f(x_i, y_j)}{\partial x} \frac{\partial f(x_i, y_j)}{\partial y} \\ \sum_{x_i, y_j \in W} \frac{\partial f(x_i, y_j)}{\partial x} \frac{\partial f(x_i, y_j)}{\partial y} & \sum_{x_i, y_j \in W} \frac{\partial^2 f(x_i, y_j)}{\partial y^2} \end{bmatrix}$$

- Usually the derivatives are computed upon a Gaussian window
- Matrix A is symmetric, positive semi-definite (non-negative eigenvalues)
- Corner information is associated with the main variation modes of the Harris matrix

Eigen-values of Harris Matrix


- Relative weight of λ_1, λ_2 , eigen-values correspond to the variation along the orthogonal directions
 - Both eigen-values are small. This means that image f is flat in the examined pixel, there are no edges or corners in this location.
 - One eigen-value is small and the second one large. The local neighborhood is ridge-shaped. Significant change of image f occurs if a small movement is made perpendicularly to the ridge.
 - Both eigen-values are rather large. A small shift in any direction causes significant change of image f . A corner is found.



(a)



(b)



(c)

Adapted from Image Processing , Analysis, and Machine Vision 3Ed, Sonka et al

Implementation issues

- Exact computation of eigen-values is not trivial
- Harris proposed a response function

$$R(A) = \det(A) - k \cdot \text{trace}^2(A)$$

- k is a heuristic coefficient, $k \in [0.04, 0.15]$


Algorithm

- 1 Filter the image with a Gaussian.
- 2 Estimate intensity gradient in two perpendicular directions for each pixel. Use successive $1D$ convolutions with the kernel approximating the derivatives
- 3 For each pixel and a given neighborhood window:
 - 1 Calculate the local structure matrix A
 - 2 Evaluate the response function $R(A)$.
 - 3 Choose the best candidates for corners by selecting a threshold on the response function $R(A)$ and perform non-maximal suppression.

Harris corner detection: remarks

- Advantages
 - Insensitivity
 - 2D shift and rotation
 - Small illumination variations
 - Small viewpoint change
 - Low computational requirements.
- Disadvantages
 - Not invariant to
 - large scale change
 - Large viewpoint changes
 - significant changes in contrast.

Harris corner detection: examples



Corners at image limits are not removed